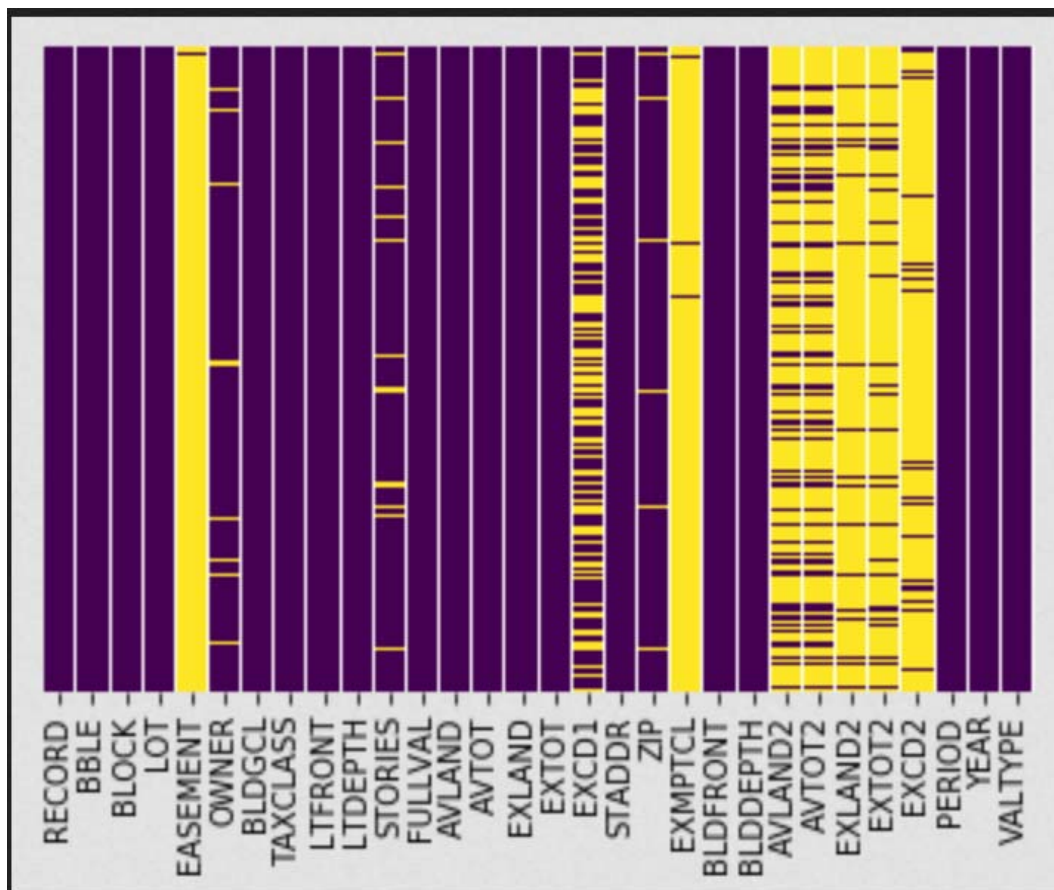
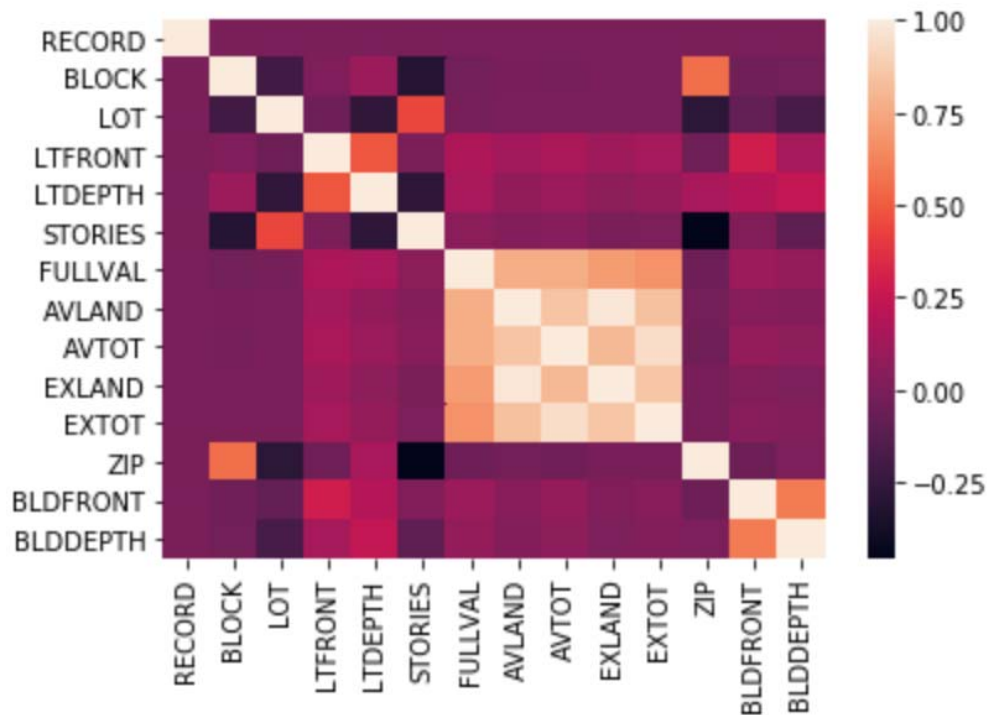


To: Professor Stephen Coggeshall
From: Alok Abhishek
Date: 01/24/2018
Subject: DSO 562: Fraud Analytics

Before starting data analysis and visualization I created a heatmap for missing values. I then removed the columns from deeper analysis because so much data is missing that interpreting few entries does not help me in looking at the bigger picture. (I've added python code for data analysis and visualization at the end the assignment)

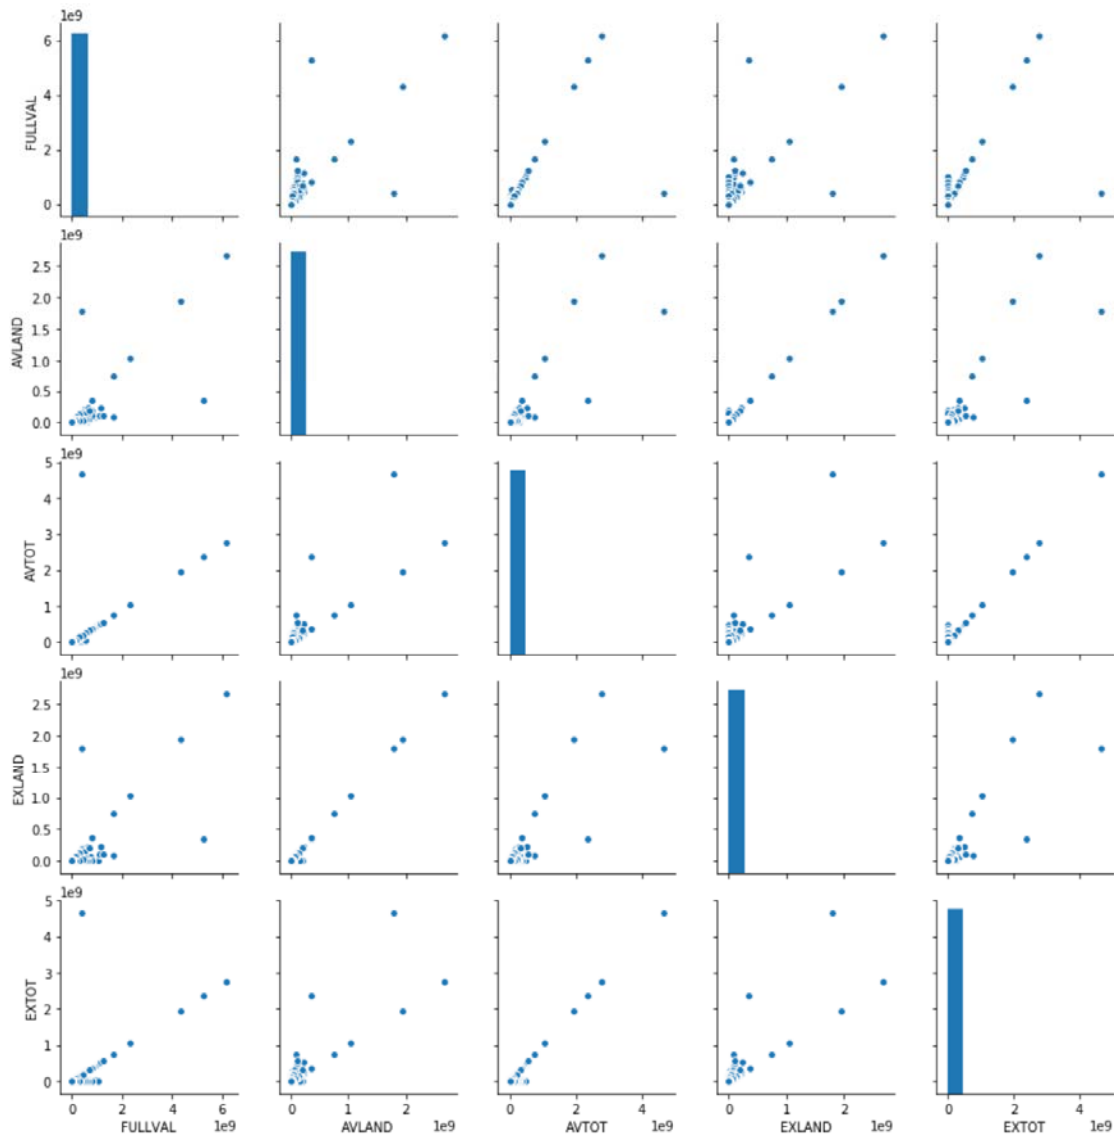


I then created heat map of correlation in between different variables to identify trends in data. I could use this relationship within data to identify outliers and potential fraud candidates.

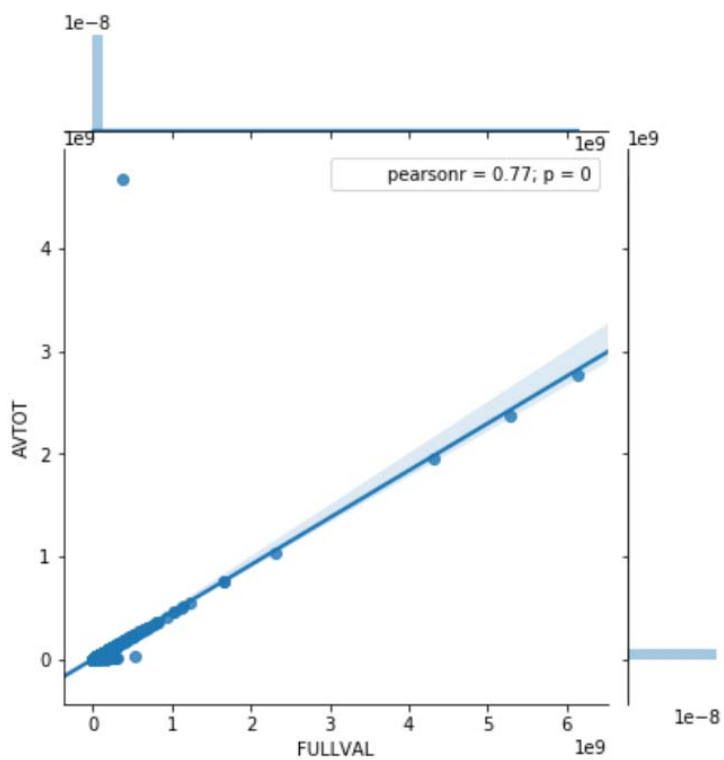





From the correlation plot we can see that there is high correlation in between Full Value, Average Land, Average Total, Ex Land, and Ex Total. I did a pair plot for these variables to identify the outliers.



Most of the variable looks well aligned across the 45 degree slope with few outliers.



Following is more in depth look at some plots which has out of pattern data points. These data points which resides outside the shaded regression region are good candidate to look at for fraud. Some of these will have valid reason and some of them could be data entry error and some may be fraudulent data.







While looking at the stories of buildings I noticed that there are several buildings with # of stories in decimals. This looks odd.

```
NYC_Property_data.groupby('STORIES').count()['RECORD']
```

```
STORIES
1.0      93606
1.1         3
1.2        33
1.3         3
1.4         2
1.5     24354
1.6     8816
1.7     5051
1.8        21
1.9        10
2.0    403318
2.1         1
2.2        40
2.3        19
2.4         2
2.5     81304
2.6        226
2.7     13543
2.8         3
2.9         1
3.0    128493
3.2        14
3.3         5
3.5     1188
3.6        11
3.7        251
4.0    38337
4.2         1
4.5        290
4.7         10
```

I also looked at lot front, lot depth, building front, and building depth. I noticed that there are a lot of properties with value zero for these fields which looks odd because how can building have 0 lot front or 0 lot depth.

```
NYC_Property_data_2.groupby('LTFRONT').count()['RECORD']
```

```
LTFRONT
0      168867
1        819
2        750
```

```
NYC_Property_data_2.groupby('LTDEPTH').count()['RECORD']
```

```
LTDEPTH
0      169888
1       126
2        79
3         1
```

```
NYC_Property_data_2.groupby('BLDFRONT').count()['RECORD']
```

```
BLDFRONT
0      224661
1         70
2         20
3         14
```


```
NYC_Property_data_2.groupby('BLDDEPTH').count()['RECORD']
```

```
BLDDEPTH
0      224699
1         52
2          9
3         90
4         60
```

Looking at the lot size in depth it seems like lot depth increases in increment of 100s.


```
tmp = NYC_Property_data[NYC_Property_data['LOT']<=1500]
tmp = tmp[tmp['LOT']>=1000]
sbrn.distplot(tmp['LOT'],bins=50)
```

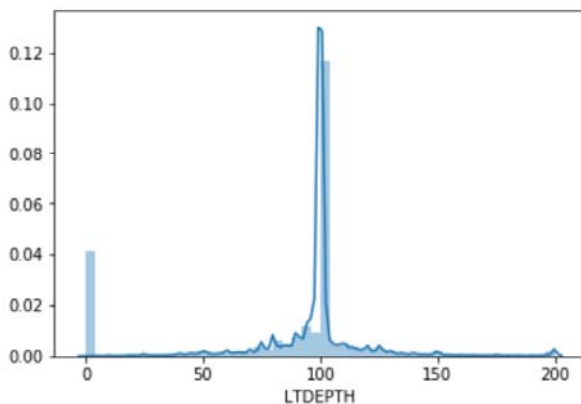
<matplotlib.axes._subplots.AxesSubplot at 0x1766618e470>



Looking at the lot depth it seems interesting that most of the lot depths are approximately 100.

```
sbrn.distplot(NYC_Property_data_2[NYC_Property_data_2['LTDEPTH']<=200]['LTDEPTH'],bins=50)
```

<matplotlib.axes._subplots.AxesSubplot at 0x2108cd91940>



```
In [2]: import numpy as np
import pandas as pd
import sklearn as sk
import seaborn as sbrn
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: NYC_Property_data = pd.read_csv('NY property 1 million.csv')
```

```
In [4]: NYC_Property_data.describe()
```

Out[4]:


	RECORD	BLOCK	LOT	LTFRONT	LTDEPTH	STORIES	F
count	1.048575e+06	1.048575e+06	1.048575e+06	1.048575e+06	1.048575e+06	996433.000000	1.048
mean	5.242880e+05	4.708867e+03	3.700924e+02	3.617425e+01	8.827643e+01	5.063363	8.804
std	3.026977e+05	3.699547e+03	8.605382e+02	7.373356e+01	7.547885e+01	8.431372	1.170
min	1.000000e+00	1.000000e+00	1.000000e+00	0.000000e+00	0.000000e+00	1.000000	0.000
25%	2.621445e+05	1.534000e+03	2.300000e+01	1.900000e+01	8.000000e+01	2.000000	3.030
50%	5.242880e+05	3.944000e+03	4.900000e+01	2.500000e+01	1.000000e+02	2.000000	4.460
75%	7.864315e+05	6.797000e+03	1.460000e+02	4.000000e+01	1.000000e+02	3.000000	6.190
max	1.048575e+06	1.635000e+04	9.978000e+03	9.999000e+03	9.999000e+03	119.000000	6.150

```
In [5]: NYC_Property_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1048575 entries, 0 to 1048574
Data columns (total 30 columns):
RECORD      1048575 non-null int64
BBLE        1048575 non-null object
BLOCK       1048575 non-null int64
LOT          1048575 non-null int64
EASEMENT     4043 non-null object
OWNER        1017492 non-null object
BLDGCL       1048575 non-null object
TAXCLASS     1048575 non-null object
LTFRONT      1048575 non-null int64
LTDEPTH      1048575 non-null int64
STORIES      996433 non-null float64
FULLVAL      1048575 non-null int64
AVLAND       1048575 non-null int64
AVTOT        1048575 non-null int64
EXLAND       1048575 non-null int64
EXTOT        1048575 non-null int64
EXCD1        622642 non-null float64
STADDR       1047934 non-null object
ZIP          1022219 non-null float64
EXMPTCL      14992 non-null object
BLDFRONT     1048575 non-null int64
BLDDEPTH     1048575 non-null int64
AVLAND2      280966 non-null float64
AVTOT2       280972 non-null float64
EXLAND2      86675 non-null float64
EXTOT2       129933 non-null float64
EXCD2        90941 non-null float64
PERIOD       1048575 non-null object
YEAR         1048575 non-null object
VALTYPE      1048575 non-null object
dtypes: float64(8), int64(12), object(10)
memory usage: 240.0+ MB
```


```
In [6]: sbrn.heatmap(NYC_Property_data.isnull(),yticklabels=False,cbar=False,cmap='viridis'
)
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x21090ac4cf8>
```



```
In [10]: sbrn.heatmap(NYC_Property_data.corr(),annot=False)
```


```
Out[10]: <matplotlib.axes._subplots.AxesSubplot at 0x17613a56f98>
```



```
In [7]: NYC_Property_data_2 = NYC_Property_data.drop(['EASEMENT', 'EXCD1','EXMPTCL','EXCD1',
    'AVLAND2', 'AVTOT2', 'EXLAND2', 'EXTOT2', 'EXCD2'], axis=1)
```

```
In [8]: sbrn.heatmap(NYC_Property_data_2.corr(),annot=False)
```

```
Out[8]: <matplotlib.axes._subplots.AxesSubplot at 0x2108d939390>
```




```
In [41]: NYC_Property_data['LOT'].describe()
```

```
Out[41]: count    1.048575e+06
mean      3.700924e+02
std       8.605382e+02
min       1.000000e+00
25%       2.300000e+01
50%       4.900000e+01
75%       1.460000e+02
max       9.978000e+03
Name: LOT, dtype: float64
```

```
In [42]: sbrn.boxplot(NYC_Property_data['LOT'])
```

```
Out[42]: <matplotlib.axes._subplots.AxesSubplot at 0x176665884e0>
```




```
In [36]: NYC_Property_data_2.groupby('LOT').count()['RECORD']
```

```
Out[36]: LOT
          1      23570
          2      6552
          3      9503
          4      8993
          5     10433
          6     11418
          7     11070
          8     10673
          9     10872
         10     10876
         11     10773
         12     11894
         13     11086
         14     11864
         15     11904
         16     11810
         17     11728
         18     11763
         19     11408
         20     12045
         21     11593
         22     11462
         23     11469
         24     11392
         25     11692
         26     11390
         27     11107
         28     11170
         29     11149
         30     11354
          ...
        9102         2
        9103         1
        9104         1
        9105         1
        9106         1
        9107         1
        9108         1
        9109         1
        9110         2
        9111         1
        9112         1
        9113         1
        9114         1
        9115         1
        9116         1
        9117         1
        9121         3
        9130         1
        9132         1
        9134         1
        9150         1
        9172         1
        9220         1
        9300         1
        9401         1
        9421         1
        9450         1
        9502         1
        9878         1
        9978         1
Name: RECORD, Length: 6366, dtype: int64
```


```
In [38]: sbrn.distplot(NYC_Property_data['LOT'],bins=50)
```

```
Out[38]: <matplotlib.axes._subplots.AxesSubplot at 0x17665f71a20>
```




```
In [40]: tmp = NYC_Property_data[NYC_Property_data['LOT']<=1500]
tmp = tmp[tmp['LOT']>=1000]
sbrn.distplot(tmp['LOT'],bins=50)
```

```
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x1766618e470>
```




```
In [37]: sbrn.distplot(NYC_Property_data[NYC_Property_data['LOT']<=200]['LOT'],bins=50,kde=T
 rue)
```

```
Out[37]: <matplotlib.axes._subplots.AxesSubplot at 0x17665fff438>
```





```
In [43]: sbrn.boxplot(x='STORIES',y='LOT',data=NYC_Property_data,palette='rainbow')
```

```
Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x17666632400>
```




```
In [44]: sbrn.boxplot(x='ZIP',y='LOT',data=NYC_Property_data,palette='rainbow')
```

```
Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x17615541320>
```



```
In [45]: sbrn.boxplot(x='ZIP',y='STORIES',data=NYC_Property_data,palette='rainbow')
```

```
Out[45]: <matplotlib.axes._subplots.AxesSubplot at 0x176157eaba8>
```




```
In [50]: NYC_Property_data['STORIES'].describe()
```

```
Out[50]: count      996433.000000  
         mean         5.063363  
         std          8.431372  
         min          1.000000  
         25%          2.000000  
         50%          2.000000  
         75%          3.000000  
         max         119.000000  
         Name: STORIES, dtype: float64
```

```
In [51]: sbrn.boxplot(NYC_Property_data['STORIES'])
```

```
Out[51]: <matplotlib.axes._subplots.AxesSubplot at 0x1761bdeb4e0>
```



```
In [58]: NYC_Property_data.groupby('STORIES').count()['RECORD']
```

Out[58]: STORIES

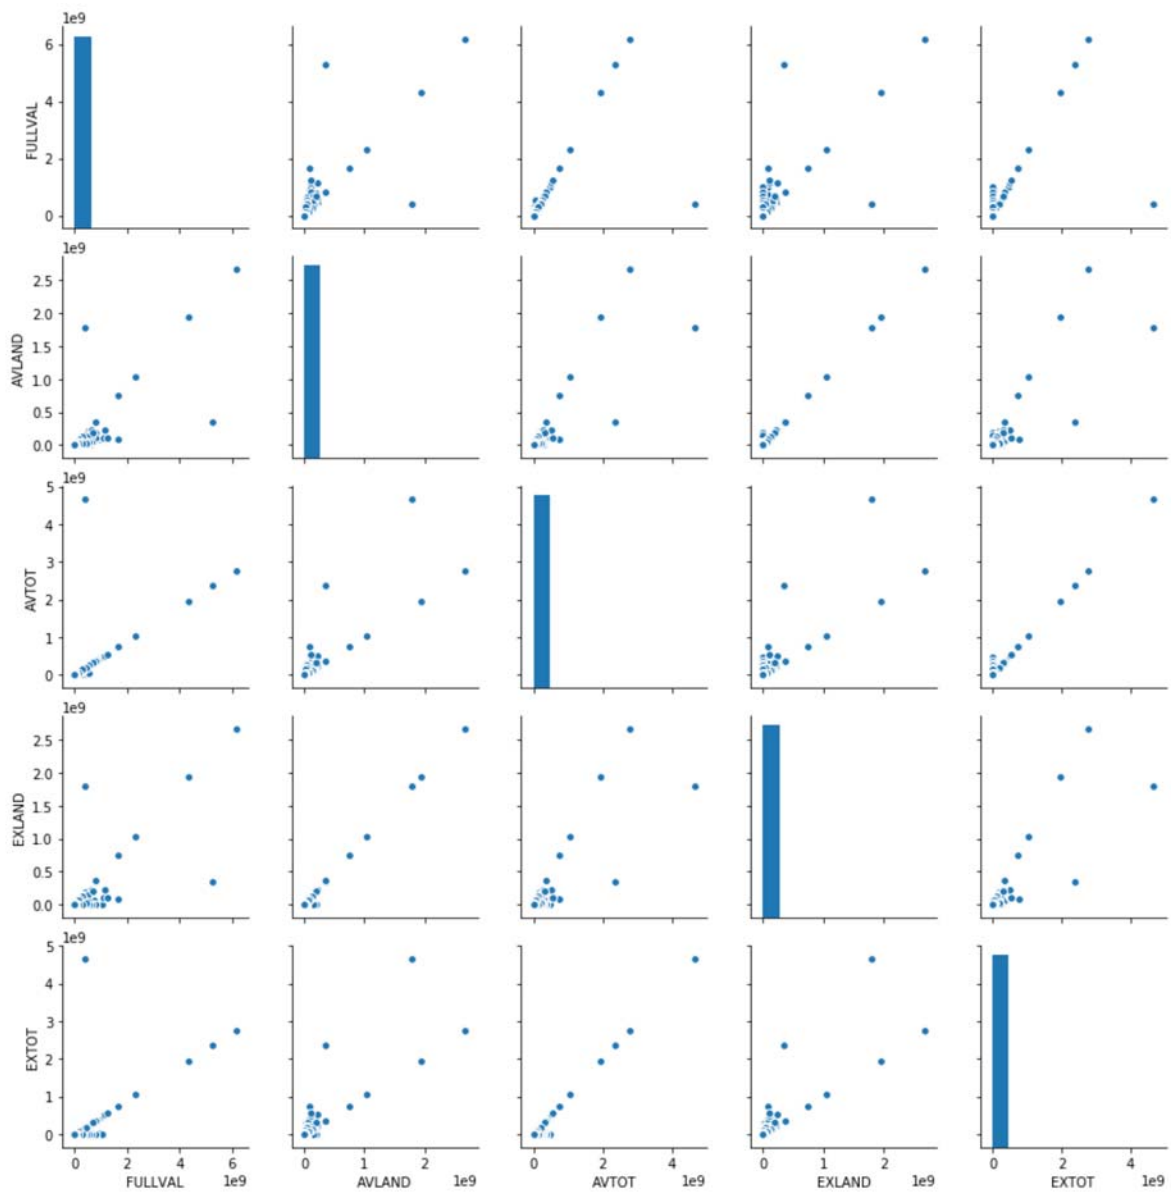
1.0	93606
1.1	3
1.2	33
1.3	3
1.4	2
1.5	24354
1.6	8816
1.7	5051
1.8	21
1.9	10
2.0	403318
2.1	1
2.2	40
2.3	19
2.4	2
2.5	81304
2.6	226
2.7	13543
2.8	3
2.9	1
3.0	128493
3.2	14
3.3	5
3.5	1188
3.6	11
3.7	251
4.0	38337
4.2	1
4.5	290
4.7	10
...	
48.0	861
49.0	472
50.0	1214
51.0	103
52.0	344
53.0	64
54.0	366
55.0	380
56.0	226
57.0	1445
58.0	253
59.0	12
60.0	561
61.0	1
62.0	3
63.0	2
65.0	72
66.0	66
67.0	242
68.0	2
70.0	849
74.0	6
75.0	31
76.0	1
78.0	1
82.0	1
85.0	1
100.0	5
114.0	1
119.0	1

Name: RECORD, Length: 111, dtype: int64

```
In [9]: NYC_Property_data_3 = NYC_Property_data_2[['FULLVAL', 'AVLAND', 'AVTOT', 'EXLAND', 'EXTOT']]
```

```
sbrn.pairplot(NYC_Property_data_3)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x2108c0d6438>
```



```
In [14]: NYC_Property_data_2.groupby('LTFRONT').count()['RECORD']
```

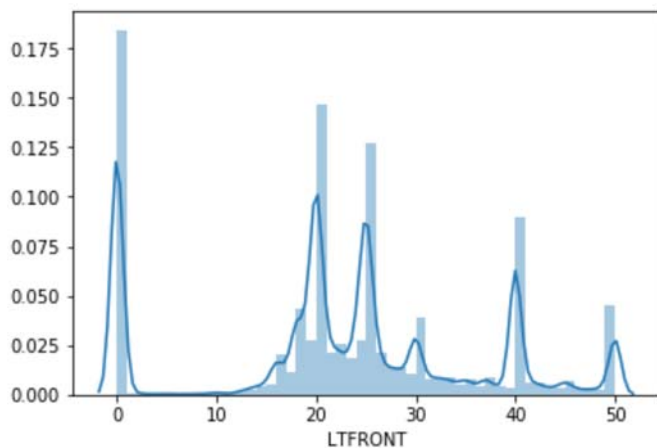
Out[14]: LTFRONT

0	168867
1	819
2	750
3	304
4	269
5	505
6	231
7	270
8	363
9	403
10	1139
11	294
12	1172
13	2327
14	4027
15	4864
16	18359
17	10372
18	40188
19	25185
20	134447
21	19319
22	23304
23	16801
24	25180
25	116301
26	19415
27	12485
28	12963
29	9249
...	
4129	1
4152	1
4171	1
4300	1
4318	1
4507	1
4644	1
4646	1
4775	1
4824	1
4910	1
4989	1
5262	1
5370	1
5380	1
5400	1
5425	1
5878	1
6078	1
6317	1
6500	1
7536	1
7653	1
8000	2
8715	2
8744	1
8821	2
9170	1
9742	1
9999	3

Name: RECORD, Length: 1277, dtype: int64

```
In [19]: sbrn.distplot(NYC_Property_data_2[NYC_Property_data_2['LTFRONT']<=50]['LTFRONT'],bins=50)
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x2108cc56ba8>
```




```
In [16]: NYC_Property_data_2.groupby('LTDEPTH').count()['RECORD']
```

Out[16]: LTDEPTH

0	169888
1	126
2	79
3	81
4	85
5	180
6	64
7	89
8	75
9	79
10	547
11	73
12	100
13	82
14	78
15	268
16	145
17	160
18	451
19	154
20	656
21	201
22	259
23	187
24	305
25	1881
26	291
27	306
28	233
29	220

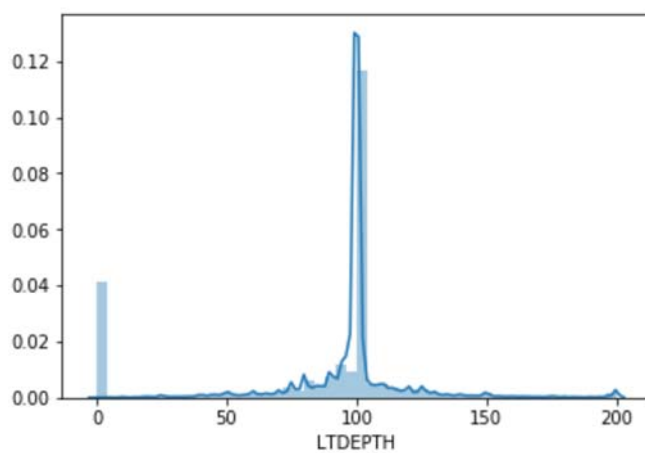
...

3700	2
3756	1
3900	1
4000	2
4050	2
4056	1
4356	1
4463	1
4471	1
4500	1
4563	1
4720	1
4770	1
4900	1
4934	1
5000	2
5100	1
5143	1
5360	2
5463	1
5853	1
5948	1
6074	1
6400	1
7055	1
7960	1
8000	1
8847	1
9619	1
9999	1

Name: RECORD, Length: 1336, dtype: int64

```
In [20]: sbrn.distplot(NYC_Property_data_2[NYC_Property_data_2['LTDEPTH']<=200]['LTDEPTH'],bins=50)
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x2108cd91940>
```



```
In [17]: NYC_Property_data_2.groupby('BLDFRONT').count()['RECORD']
```

Out[17]: BLDFRONT

0	224661
1	70
2	20
3	14
4	14
5	45
6	33
7	23
8	209
9	197
10	1385
11	143
12	2076
13	4854
14	15792
15	16013
16	73671
17	23821
18	76808
19	33073
20	193812
21	32593
22	53227
23	16036
24	32472
25	61770
26	28443
27	15112
28	9035
29	3841

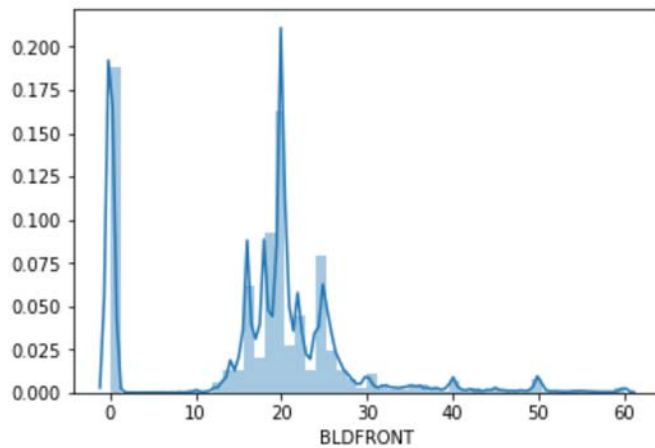
...

900	1
911	1
961	1
982	1
1000	2
1102	1
1160	1
1169	1
1225	1
1227	1
1280	1
1362	1
1394	1
1812	1
1844	2
1925	1
1943	1
2025	1
2030	1
2500	1
3100	1
3285	1
4017	1
4149	1
5518	1
5614	1
6020	1
6414	1
7538	1
7575	1

Name: RECORD, Length: 610, dtype: int64

```
In [21]: sbrn.distplot(NYC_Property_data_2[NYC_Property_data_2['BLDFRONT']<=60]['BLDFRONT'],  
bins=50)
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x2108ce11ba8>
```



```
In [18]: NYC_Property_data_2.groupby('BLDDEPTH').count()['RECORD']
```

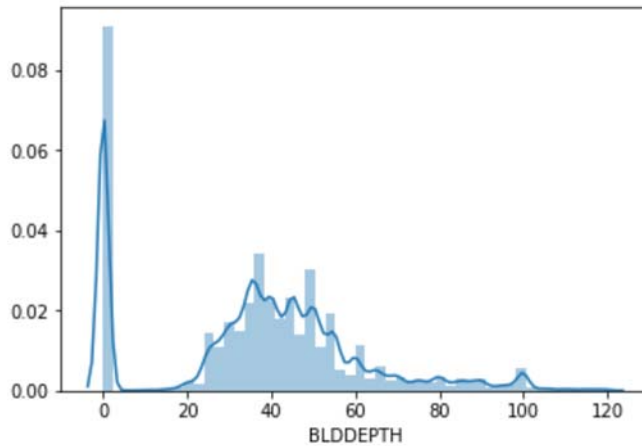
Out[18]: BLDDEPTH

0	224699
1	52
2	9
3	90
4	60
5	51
6	48
7	13
8	53
9	14
10	535
11	22
12	174
13	54
14	133
15	673
16	416
17	213
18	1345
19	339
20	4305
21	896
22	3020
23	1459
24	9892
25	11491
26	14512
27	7243
28	19709
29	4605
...	
980	3
992	2
999	2
1000	2
1007	1
1075	1
1131	1
1150	1
1175	1
1222	1
1300	2
1375	1
1399	1
1971	1
1980	1
2023	1
2436	1
3104	1
3390	1
4500	1
4600	1
5000	1
5020	1
5600	1
5641	1
6308	1
7360	1
8500	1
9388	1
9393	1

Name: RECORD, Length: 620, dtype: int64



```
In [22]: sbrn.distplot(NYC_Property_data_2[NYC_Property_data_2['BLDDEPTH'] <= 120]['BLDDEPTH'], bins=50)
```

```
Out[22]: <matplotlib.axes._subplots.AxesSubplot at 0x2108e21ee48>
```



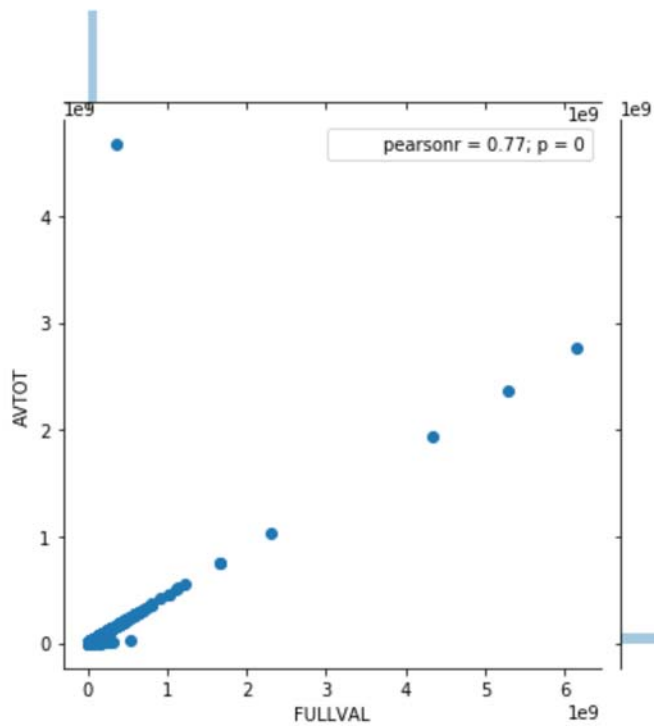
```
In [25]: sbrn.jointplot(x='FULLVAL', y='AVLAND', data=NYC_Property_data_3)
```

```
Out[25]: <seaborn.axisgrid.JointGrid at 0x2108e483b00>
```



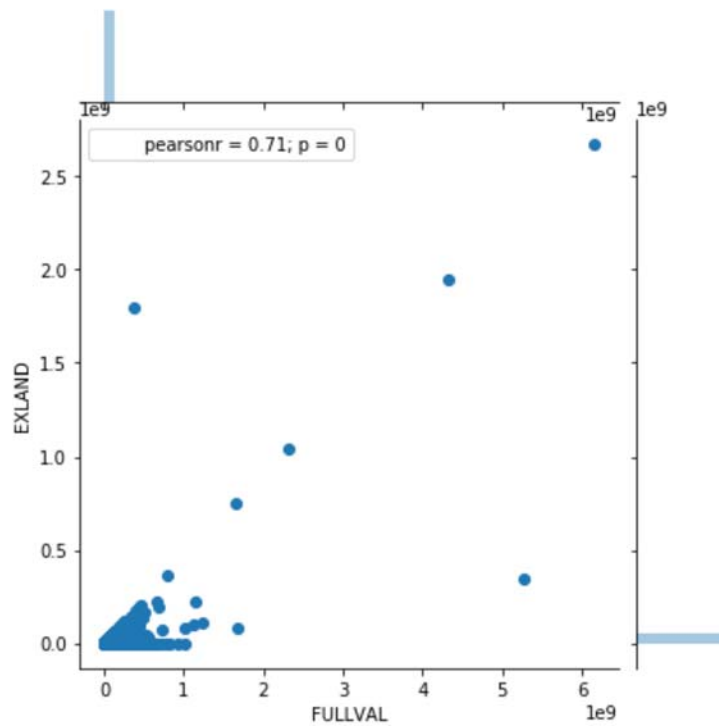
```
In [26]: sbrn.jointplot(x='FULLVAL',y='AVTOT',data=NYC_Property_data_3)
```

```
Out[26]: <seaborn.axisgrid.JointGrid at 0x2108e4961d0>
```



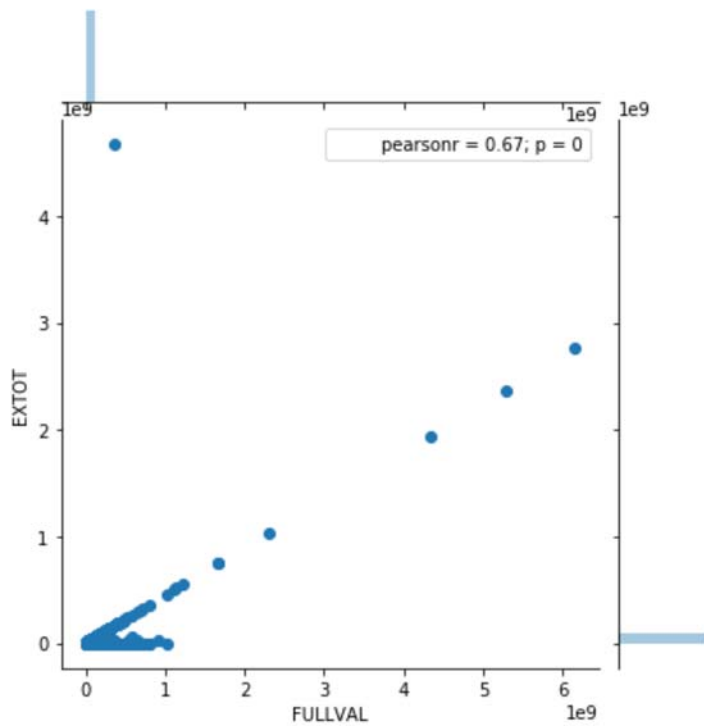
```
In [27]: sbrn.jointplot(x='FULLVAL',y='EXLAND',data=NYC_Property_data_3)
```

```
Out[27]: <seaborn.axisgrid.JointGrid at 0x2108e645898>
```



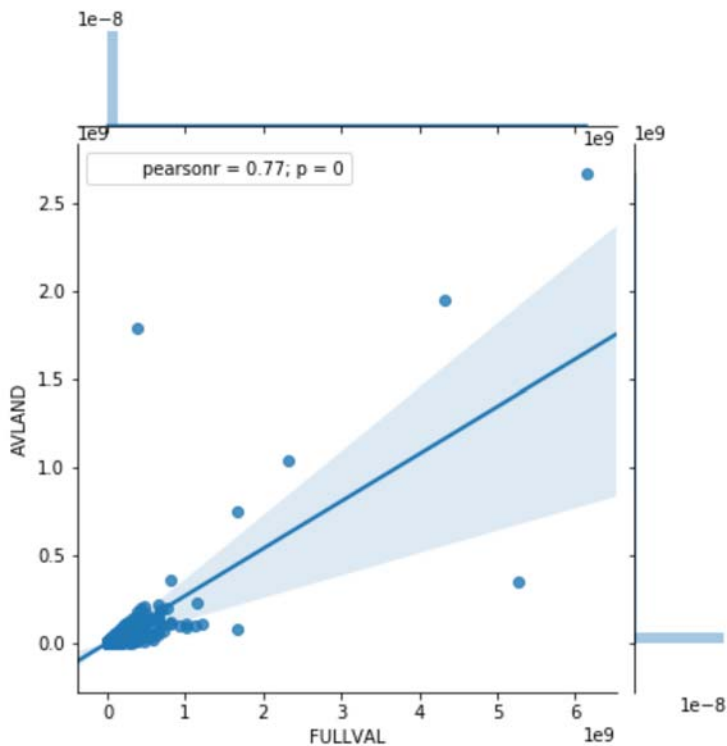
```
In [28]: sbrn.jointplot(x='FULLVAL',y='EXTOT',data=NYC_Property_data_3)
```

```
Out[28]: <seaborn.axisgrid.JointGrid at 0x2108e84f208>
```



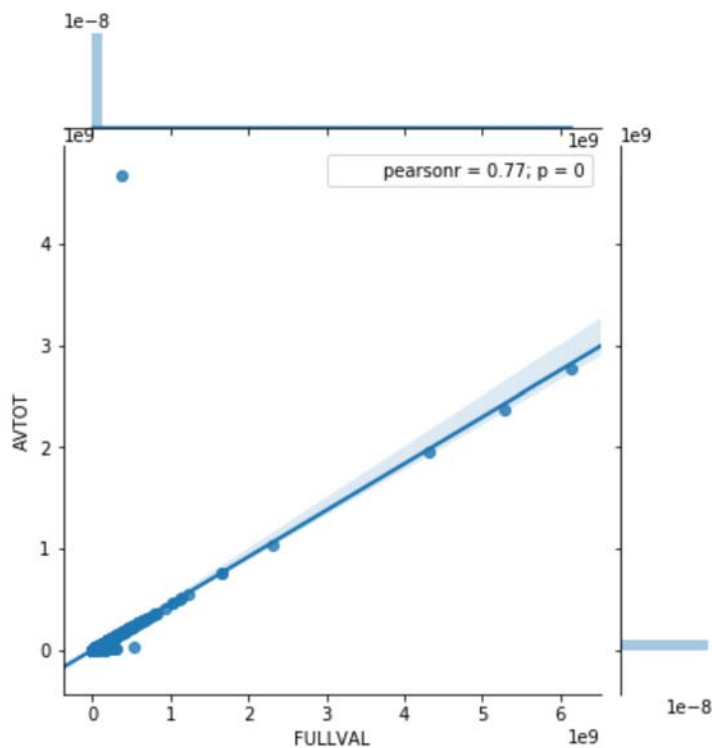
```
In [29]: sbrn.jointplot("FULLVAL", "AVLAND", data=NYC_Property_data_3, kind="reg")
```

```
Out[29]: <seaborn.axisgrid.JointGrid at 0x2108eb0fe48>
```




```
In [30]: sbrn.jointplot("FULLVAL", "AVTOT", data=NYC_Property_data_3, kind="reg")
```

```
Out[30]: <seaborn.axisgrid.JointGrid at 0x2108eb0f3c8>
```



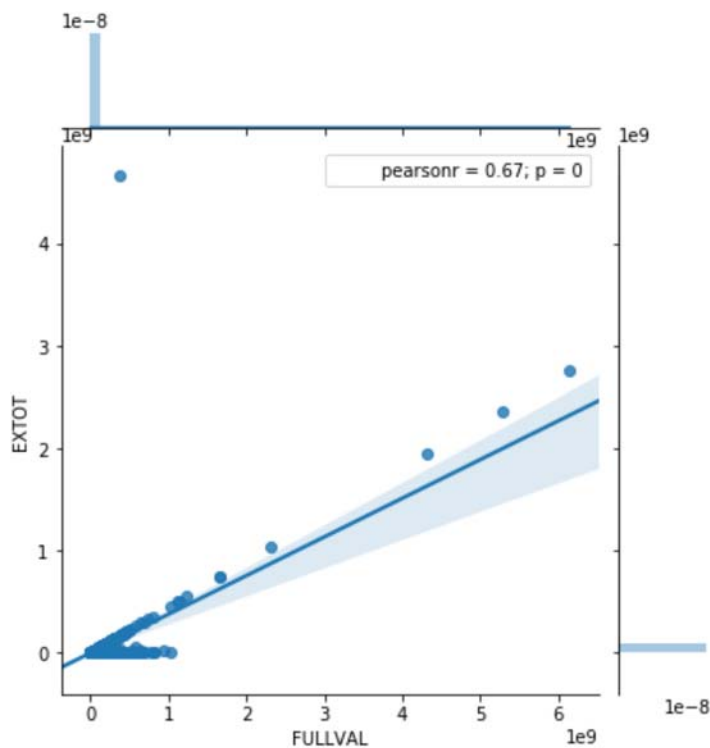
```
In [31]: sbrn.jointplot("FULLVAL", "EXLAND", data=NYC_Property_data_3, kind="reg")
```

```
Out[31]: <seaborn.axisgrid.JointGrid at 0x2108ff3d710>
```



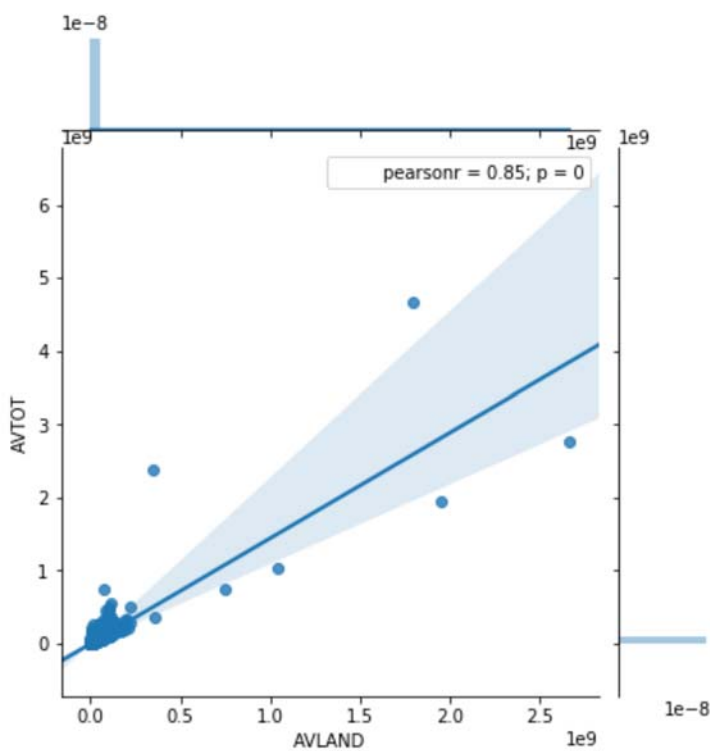
```
In [32]: sbrn.jointplot("FULLVAL", "EXTOT", data=NYC_Property_data_3, kind="reg")
```

```
Out[32]: <seaborn.axisgrid.JointGrid at 0x21090a532b0>
```



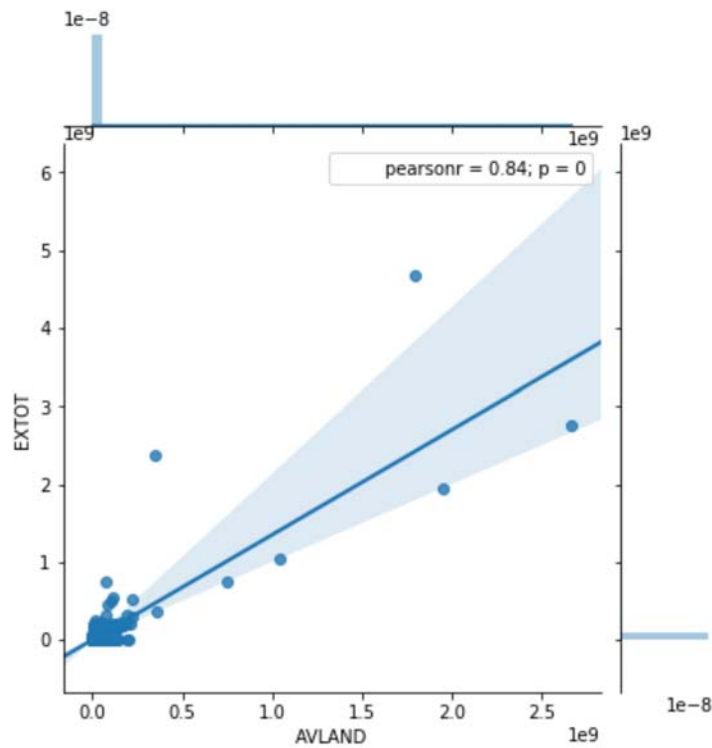
```
In [37]: sbrn.jointplot("AVLAND", "AVTOT", data=NYC_Property_data_3, kind="reg")
```

```
Out[37]: <seaborn.axisgrid.JointGrid at 0x210a02e2978>
```



```
In [38]: sbrn.jointplot("AVLAND", "EXTOT", data=NYC_Property_data_3, kind="reg")
```

```
Out[38]: <seaborn.axisgrid.JointGrid at 0x210a2a08da0>
```



```
In [39]: sbrn.jointplot("AVLAND", "EXLAND", data=NYC_Property_data_3, kind="reg")
```

```
Out[39]: <seaborn.axisgrid.JointGrid at 0x210a2a33898>
```

